

Cherry Blossom Prediction Competition

Spencer Retcher

Importing Tidyverse

```
library(tidyverse)
library(caret)
library(openmeteo)
library(readr)
library(kernlab)
```

Measuring Bloom Dates for New York City

The peak bloom dates for Yoshino cherry trees from 2019 - 2023 were determined by using data from the *USA National Phenology Network* and by following this method:

1. The bloom date for each tree was measured as the average date between the day the volunteer first noticed open flowers and the number of days since the prior observation.
 - i) If a volunteer noticed open flowers on a certain day, that does not mean that the bloom date was that day. The bloom date could have taken place anytime between these two dates, so this average attempts to minimize potential volunteer error.
2. The peak bloom dates for New York City in 2019-2023 were the earliest bloom dates observed among Yoshino cherry trees during those years.

```
nyc <- read_csv("data/USA-NPN_individual_phenometrics_data.csv") |>
  filter(Site_ID == 32789, Species_ID == 228) |>
  mutate(NumDays_Since_Prior_No = if_else(NumDays_Since_Prior_No == -9999,
                                           0,
                                           NumDays_Since_Prior_No))

nyc_data <- nyc |>
```

```

mutate(
  year = First_Yes_Year,
  bloom_date = as.Date(paste(First_Yes_Year, First_Yes_Month, First_Yes_Day, sep = "-"))
  bloom_doy = floor(First_Yes_DOY - (NumDays_Since_Prior_No / 2))
) |>
group_by(year) |>
summarise(bloom_date = min(bloom_date),
          bloom_doy = min(bloom_doy),
          .groups = "drop") |>
mutate(
  location = 'newyorkcity',
  lat = 40.73082,
  long = -73.99733,
  alt = 5,
) |>
select(location, lat, long, alt, year, bloom_date, bloom_doy)

```

nyc_data

```

# A tibble: 4 x 7
  location    lat long  alt year bloom_date bloom_doy
  <chr>      <dbl> <dbl> <dbl> <dbl> <date>      <dbl>
1 newyorkcity 40.7 -74.0    5 2019 2019-04-05    95
2 newyorkcity 40.7 -74.0    5 2021 2021-04-04    94
3 newyorkcity 40.7 -74.0    5 2022 2022-03-26    85
4 newyorkcity 40.7 -74.0    5 2023 2023-03-26    85

```

Historical Bloom Dates

The tibble `historical_bloom_dates` contains past peak bloom dates and information about the five locations of interest:

- Kyoto (Japan)
- Liestal-Weideli (Switzerland)
- Washington, D.C. (USA)
- Vancouver, BC (Canada)
- New York City, NY (USA)

```

historical_bloom_dates <- tibble(read_csv("data/washingtondc.csv") |>
  bind_rows(read_csv("data/liestal.csv"))) |>

```

```
bind_rows(read_csv("data/kyoto.csv")) |>
bind_rows(read_csv("data/vancouver.csv")) |>
bind_rows(nyc_data))
```

```
historical_bloom_dates
```

```
# A tibble: 1,074 x 7
```

	location	lat	long	alt	year	bloom_date	bloom_doy
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<date>	<dbl>
1	washingtondc	38.9	-77.0	0	1921	1921-03-20	79
2	washingtondc	38.9	-77.0	0	1922	1922-04-07	97
3	washingtondc	38.9	-77.0	0	1923	1923-04-09	99
4	washingtondc	38.9	-77.0	0	1924	1924-04-13	104
5	washingtondc	38.9	-77.0	0	1925	1925-03-27	86
6	washingtondc	38.9	-77.0	0	1926	1926-04-11	101
7	washingtondc	38.9	-77.0	0	1927	1927-03-20	79
8	washingtondc	38.9	-77.0	0	1928	1928-04-08	99
9	washingtondc	38.9	-77.0	0	1929	1929-03-31	90
10	washingtondc	38.9	-77.0	0	1930	1930-04-01	91

```
# i 1,064 more rows
```

Open-Meteo API

Data for this project was collected from the *Open-Meteo API*, which can be accessed from R by using the library `openmeteo` (API key not required). Using this API, we will retrieve hourly and daily weather conditions from 1940 to the present day for each of our five locations. Later on, we will forecast the same weather conditions 16 days into the future for each of our locations, which we will use to help us make our 2024 peak bloom predictions.

Note: The API can be slow and allows around 10-15 requests per day. I will convert most of these requests to csv files and save them to the folders `data/hourly_data`, `data/daily_data`, and `data/forecast_data` to make the report easier to reproduce.

```
# Used to get daily weather data for a location
get_daily_data <- function(lat, long, date_min, date_max, elements) {
  weather_history(
    location = c(lat, long),
    start = date_min,
    end = date_max,
    daily = elements
```

```

    )
  }

# Used to get hourly weather data for a location
get_hourly_data <- function(lat, long, date_min, date_max, elements) {
  weather_history(
    location = c(lat, long),
    start = date_min,
    end = date_max,
    hourly = elements
  )
}

```

Accumulative Growing Degree Days

Growing degree day A measurement of heat accumulation which can be used to predict the development rates of plants.

`add_gdd()` calculates the growing degree day for each day by using the

1. Daily maximum temperature
2. Daily minimum temperature
3. Base Temperature of 10°C

$$GDD = \frac{(T_{max} + T_{min})}{2} - 10^{\circ}C$$

If the average daily temperature is less than the base temperature, the growing degree day for that day is zero.

`get_agdd()` filters the daily weather data for a particular year and location and sums up the growing degree days for that year between January 1st to the `bloom_date`.

```

add_gdd <- function(daily_temp_data) {
  daily_temp_data |>
    mutate(
      year = year(date),
      growing_degree_day =
        ((daily_temperature_2m_max + daily_temperature_2m_min) / 2) - 10
    ) |>
    mutate(growing_degree_day = case_when(growing_degree_day < 0 ~ 0,
                                           .default = growing_degree_day)) |>

```

```

    relocate(year, .after = date)
  }

get_agdd <- function(place, year, bloom_date) {

  data <- daily_data %>%
    filter(location == place)

  date_min = as.Date(paste(year, "-01-01", sep = ""), format = "%Y-%m-%d")

  check_date_range <- data %>%
    filter(date >= date_min & date <= bloom_date)

  if (nrow(check_date_range) == 0) {
    return(NA)
  }

  check_date_range |>
    group_by(year) |>
    summarise(agdd = sum(growing_degree_day, na.rm = TRUE)) |>
    pull(agdd)
}

```

Total Sunshine Duration

Sunshine duration The number of seconds of sunshine per day which is determined by calculating direct normalized irradiance exceeding 120 W/m².

`sunshine_to_hours()` converts the sunshine duration of each day into hours.

`get_sunshine()` filters the daily weather data for a particular year and location and sums up the number of hours of sunshine duration for that year between January 1st to the `bloom_date`.

```

sunshine_to_hours <- function(sunshine_data) {
  sunshine_data |>
    mutate(daily_sunshine_duration = daily_sunshine_duration / 3600)
}

```

```

get_sunshine <- function(place, year, bloom_date) {

  data <- daily_data %>%
    filter(location == place)

  date_min = as.Date(paste(year, "-01-01", sep = ""), format = "%Y-%m-%d")

  check_date_range <- data %>%
    filter(date >= date_min & date <= bloom_date)

  if (nrow(check_date_range) == 0) {
    return(NA)
  }

  check_date_range |>
    group_by(year) |>
    summarise(sunshine_duration = sum(daily_sunshine_duration, na.rm = TRUE)) |>
    pull(sunshine_duration)
}

```

Total Precipitation

Precipitation The amount of rain, showers, and snowfall for a day measured in millimeters.

`get_precip()` filters the daily weather data for a particular year and location and sums up the amount of precipitation for that year between January 1st to the `bloom_date`.

```

get_precip <- function(place, year, bloom_date) {

  data <- daily_data %>%
    filter(location == place)

  date_min = as.Date(paste(year, "-01-01", sep = ""), format = "%Y-%m-%d")

  check_date_range <- data %>%
    filter(date >= date_min & date <= bloom_date)

  if (nrow(check_date_range) == 0) {
    return(NA)
  }
}

```

```

    check_date_range |>
      group_by(year) |>
      summarise(total_precipitation = sum(daily_precipitation_sum, na.rm = TRUE)) |>
      pull(total_precipitation)
  }

```

Chill Hours

Chill hours The number of hours in the winter that a plant spends exposed to certain temperatures.

For this project, chill hours is the number of hours between 0°C to 7°C from November to February.

`get_chills()` filters the hourly weather data for a particular year and location and sums up the number of chill hours in the winter for that year between November to February.

```

add_chill <- function(data) {
  data |>
    mutate(
      chill_hour = case_when(
        hourly_temperature_2m >= 0 & hourly_temperature_2m <= 7 ~ 1,
        .default = 0
      ),
      date = as.Date(datetime),
      year = as.integer(format(datetime, "%Y")),
      month = as.integer(strftime(date, '%m')) %% 12,
      # make December "0"
      year = if_else(month == 0 | month == 11, year + 1L, year)
    ) |>
    filter(month %in% c(11, 0, 1, 2)) |>
    group_by(year, location) |>
    summarize(chill_hours = sum(chill_hour))
}

```

```

get_chill <- function(place, chill_year) {

```

```

data <- hourly_data %>%
  filter(location == place)

check_date_range <- data %>%
  filter(year == chill_year)

if (nrow(check_date_range) == 0) {
  return(NA)
}

data |>
  filter(year == chill_year) |>
  pull(chill_hours)
}

```

Forecasting Weather Data

The *Open-Meteo API* can forecast the daily weather conditions stated above 16 days into the future. These forecasts were downloaded for each location from <https://open-meteo.com/> and stored in csv files.

We do not need to forecast hourly temperature data to calculate this winter's chill hours since the competition ends on the last day of winter and we already have a good estimate of the chill hours.

```

washingtondc_forecast_daily_data <- read_csv("data/forecast_data/open-meteo-38.89N77.03W1m.csv")
rename(date = time,
        daily_temperature_2m_max = `temperature_2m_max (°C)`,
        daily_temperature_2m_min = `temperature_2m_min (°C)`,
        daily_sunshine_duration = `sunshine_duration (s)`,
        daily_precipitation_sum = `precipitation_sum (mm)` ) |>
mutate(location = "washingtondc")

lietal_forecast_daily_data <- read_csv("data/forecast_data/open-meteo-47.48N7.74E344m.csv")
rename(date = time,
        daily_temperature_2m_max = `temperature_2m_max (°C)`,
        daily_temperature_2m_min = `temperature_2m_min (°C)`,
        daily_sunshine_duration = `sunshine_duration (s)`,
        daily_precipitation_sum = `precipitation_sum (mm)` ) |>

```



```

mutate(location = "liestal")

kyoto_forecast_daily_data <- read_csv("data/forecast_data/open-meteo-35.00N135.69E31m.csv")
  rename(date = time,
         daily_temperature_2m_max = `temperature_2m_max (°C)`,
         daily_temperature_2m_min = `temperature_2m_min (°C)`,
         daily_sunshine_duration = `sunshine_duration (s)`,
         daily_precipitation_sum = `precipitation_sum (mm)`) |>
mutate(location = "kyoto")

vancouver_forecast_daily_data <- read_csv("data/forecast_data/open-meteo-49.23N123.18W28m.csv")
  rename(date = time,
         daily_temperature_2m_max = `temperature_2m_max (°C)`,
         daily_temperature_2m_min = `temperature_2m_min (°C)`,
         daily_sunshine_duration = `sunshine_duration (s)`,
         daily_precipitation_sum = `precipitation_sum (mm)`) |>
mutate(location = "vancouver")

newyorkcity_forecast_daily_data <- read_csv("data/forecast_data/open-meteo-40.74N73.98W14m.csv")
  rename(date = time,
         daily_temperature_2m_max = `temperature_2m_max (°C)`,
         daily_temperature_2m_min = `temperature_2m_min (°C)`,
         daily_sunshine_duration = `sunshine_duration (s)`,
         daily_precipitation_sum = `precipitation_sum (mm)`) |>
mutate(location = "newyorkcity")

bind_rows(washingtondc_forecast_daily_data,
          liestal_forecast_daily_data,
          kyoto_forecast_daily_data,
          vancouver_forecast_daily_data,
          newyorkcity_forecast_daily_data)

# A tibble: 80 x 6
  date           daily_temperature_2m_max daily_temperature_2m_min
  <date>                <dbl>                <dbl>
1 2024-02-29             14.6                -0.3
2 2024-03-01             11.1                 0

```

```

3 2024-03-02          9.4          6.5
4 2024-03-03         16.5          6.4
5 2024-03-04         16.8         10.2
6 2024-03-05         13.9          8.3
7 2024-03-06         14.5          10
8 2024-03-07         20          12.7
9 2024-03-08         14.4          6.8
10 2024-03-09        10.9          7.3
# i 70 more rows
# i 3 more variables: daily_sunshine_duration <dbl>,
#   daily_precipitation_sum <dbl>, location <chr>

```

Combining and Transforming Daily Weather Data and Forecast Data

1. The *Open-Meteo API* was used to retrieve the daily maximum temperature, daily minimum temperature, daily sunshine duration, and daily precipitation for all days from January 1, 1940 to Yesterday for each location.
2. All of the daily weather data was joined with the 16 day weather forecast for each location and placed into `daily_data`.
 - i) This means if we submit our project on the last day of the competition, we will have the daily weather data from January 1, 1940 to March 15, 2024.
3. `daily_data` is then transformed.
 - i) The sunshine duration for each day is converted to hours.
 - ii) The growing degree days are added for each day.

```

# washingtondc_daily_data <- get_daily_data(lat = 38.8853,
#                                           long = -77.0386,
#                                           date_min = "1940-01-01",
#                                           date_max = Sys.Date() - 1,
#                                           elements = c("temperature_2m_max", "temperature_2m_min", "precipitation_sum", "sunshine_duration"),
#                                           mutate(location = "washingtondc"))
# write.csv(washingtondc_daily_data, file = "washingtondc_daily_data.csv", row.names = TRUE)

washingtondc_daily_data <- read_csv("data/daily_data/washingtondc_daily_data.csv")

# liestal_daily_data <- get_daily_data(lat = 47.4814,
#                                       long = 7.730519 ,

```

```

#                                     date_min = "1940-01-01",
#                                     date_max = Sys.Date() - 1,
#                                     elements = c("temperature_2m_max", "temperature_2m_min"),
#                                     mutate(location = "liestal")
# write.csv(liestal_daily_data, file = "liestal_daily_data.csv", row.names = TRUE)

liestal_daily_data <- read_csv("data/daily_data/liestal_daily_data.csv")

# kyoto_daily_data <- get_daily_data(lat = 35.0120,
#                                     long = 135.6761,
#                                     date_min = "1940-01-01",
#                                     date_max = Sys.Date() - 1,
#                                     elements = c("temperature_2m_max", "temperature_2m_min"),
#                                     mutate(location = "kyoto")
# write.csv(kyoto_daily_data, file = "kyoto_daily_data.csv", row.names = TRUE)

kyoto_daily_data <- read_csv("data/daily_data/kyoto_daily_data.csv")

#
# vancouver_daily_data <- get_daily_data(lat = 49.2237,
#                                     long = -123.1636,
#                                     date_min = "1940-01-01",
#                                     date_max = Sys.Date() - 1,
#                                     elements = c("temperature_2m_max", "temperature_2m_min"),
#                                     mutate(location = "vancouver")
# write.csv(vancouver_daily_data, file = "vancouver_daily_data.csv", row.names = TRUE)

vancouver_daily_data <- read_csv("data/daily_data/vancouver_daily_data.csv")

#
# newyorkcity_daily_data <- get_daily_data(lat = 40.73040 ,
#                                     long = -73.99809,
#                                     date_min = "1940-01-01",
#                                     date_max = Sys.Date() - 1,
#                                     elements = c("temperature_2m_max", "temperature_2m_min"),
#                                     mutate(location = "newyorkcity")

# write.csv(newyorkcity_daily_data, file = "newyorkcity_daily_data.csv", row.names = TRUE)

newyorkcity_daily_data <- read_csv("data/daily_data/newyorkcity_daily_data.csv")

```

```
daily_data = bind_rows(washingtondc_daily_data, liestal_daily_data, kyoto_daily_data, vancouver_daily_data) |>
  add_gdd() |>
  sunshine_to_hours()

daily_data
```

```
# A tibble: 153,775 x 9
  ...1 date      year daily_temperature_2m_max daily_temperature_2m_min
  <dbl> <date>    <dbl>          <dbl>          <dbl>
1     1 1940-01-01  1940          -2.3          -7.6
2     2 1940-01-02  1940          -3.3          -8.1
3     3 1940-01-03  1940           -2          -8.8
4     4 1940-01-04  1940         -0.2          -7.5
5     5 1940-01-05  1940           6.3          -7.3
6     6 1940-01-06  1940          -1.2          -8.1
7     7 1940-01-07  1940          -0.8          -7.9
8     8 1940-01-08  1940           2.1          -2.6
9     9 1940-01-09  1940           0.1          -5.7
10    10 1940-01-10  1940           1.5           -7
# i 153,765 more rows
# i 4 more variables: daily_sunshine_duration <dbl>,
#   daily_precipitation_sum <dbl>, location <chr>, growing_degree_day <dbl>
```

Hourly Weather Data

The *Open-Meteo API* was used to retrieve the hourly temperatures for all days from January 1, 1940 to Yesterday for each location.

These hourly temperatures were joined together, placed into `hourly_data`, and then the chill hours for each hour were calculated.

```
#
# washingtondc_hourly_data <- get_hourly_data(lat = 38.8853,
#                                           long = -77.0386,
#                                           date_min = "1940-01-01",
#                                           date_max = Sys.Date() - 1,
#                                           elements = c("temperature_2m")) %>%
#                                           mutate(location = "washingtondc")
# write.csv(washingtondc_hourly_data, file = "washingtondc_hourly_data.csv", row.names = T)
```

```

washingtondc_hourly_data <- read_csv("data/hourly_data/washingtondc_hourly_data.csv")

#
# liestal_hourly_data <- get_hourly_data(lat = 47.4814,
#                                       long = 7.730519 ,
#                                       date_min = "1940-01-01",
#                                       date_max = Sys.Date() - 1,
#                                       elements = c("temperature_2m")) %>%
#                                       mutate(location = "liestal")

# write.csv(liestal_hourly_data, file = "liestal_hourly_data.csv", row.names = TRUE)

liestal_hourly_data <- read_csv("data/hourly_data/liestal_hourly_data.csv")

#
# kyoto_hourly_data <- get_hourly_data(lat = 35.0120,
#                                       long = 135.6761,
#                                       date_min = "1940-01-01",
#                                       date_max = Sys.Date() - 1,
#                                       elements = c("temperature_2m" )) %>%
#                                       mutate(location = "kyoto")
# write.csv(kyoto_hourly_data, file = "kyoto_hourly_data.csv", row.names = TRUE)

kyoto_hourly_data <- read_csv("data/hourly_data/kyoto_hourly_data.csv")

#
# vancouver_hourly_data <- get_hourly_data(lat = 49.2237 ,
#                                       long = -123.1636 ,
#                                       date_min = "1940-01-01",
#                                       date_max = Sys.Date() - 1,
#                                       elements = c("temperature_2m" )) %>%
#                                       mutate(location = "vancouver")
# write.csv(vancouver_hourly_data, file = "vancouver_hourly_data.csv", row.names = TRUE)

vancouver_hourly_data<- read_csv("data/hourly_data/vancouver_hourly_data.csv")

#
# newyorkcity_hourly_data <- get_hourly_data(lat = 40.73040 ,
#                                       long = -73.99809,
#                                       date_min = "1940-01-01",

```

```

#                                     date_max = Sys.Date() - 1,
#                                     elements = c("temperature_2m" )) %>%
#                                     mutate(location = "newyorkcity")
# write.csv(newyorkcity_hourly_data, file = "newyorkcity_hourly_data.csv", row.names = TRUE)

newyorkcity_hourly_data <- read_csv("data/hourly_data/newyorkcity_hourly_data.csv")

hourly_data = bind_rows(washingtondc_hourly_data,liestal_hourly_data,kyoto_hourly_data,vancouver_hourly_data,
                        add_chill())

```

`summarise()` has grouped output by 'year'. You can override using the
 ` .groups ` argument.

```
hourly_data
```

```

# A tibble: 425 x 3
# Groups:   year [85]
   year location    chill_hours
  <int> <chr>         <dbl>
1  1940 kyoto          785
2  1940 liestal        444
3  1940 newyorkcity    417
4  1940 vancouver    1178
5  1940 washingtondc   520
6  1941 kyoto        1380
7  1941 liestal      1283
8  1941 newyorkcity   1370
9  1941 vancouver    2272
10 1941 washingtondc  1422
# i 415 more rows

```

Creating the Complete Dataset

We have all the daily and hourly weather data for each of our locations from January 1, 1940 to March 15, 2024.

We then use `rowwise()` and `mutate()` to essentially take each row of data in `historical_bloom_dates` and compute the aggregations for that year.

For each year and location in `historical_bloom_dates`, we calculate:

1. the number of chill hours between November and February.
2. the number of sunshine hours between January 1st to the bloom_date.
3. the number of growing degree days between January 1st to the bloom_date.
4. the amount of precipitation between January 1st to the bloom_date.

```
complete_dataset <- historical_bloom_dates %>%
  filter(year >=1940) |>
  rowwise() |>
  mutate(chill_hours = get_chill(location,year),
         accumulative_growing_degree_days = get_agdd(location, year, bloom_date),
         total_sunshine_duration = get_sunshine(location,year,bloom_date),
         total_precipitation = get_precip(location, year, bloom_date)
  )
```

```
complete_dataset
```

```
# A tibble: 257 x 11
```

```
# Rowwise:
```

	location	lat	long	alt	year	bloom_date	bloom_doy	chill_hours
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<date>	<dbl>	<dbl>
1	washingtondc	38.9	-77.0	0	1940	1940-04-13	104	520
2	washingtondc	38.9	-77.0	0	1941	1941-04-12	102	1422
3	washingtondc	38.9	-77.0	0	1942	1942-04-05	95	1321
4	washingtondc	38.9	-77.0	0	1943	1943-04-04	94	1384
5	washingtondc	38.9	-77.0	0	1944	1944-04-09	100	1419
6	washingtondc	38.9	-77.0	0	1945	1945-03-20	79	1381
7	washingtondc	38.9	-77.0	0	1946	1946-03-23	82	1270
8	washingtondc	38.9	-77.0	0	1947	1947-04-12	102	1246
9	washingtondc	38.9	-77.0	0	1948	1948-03-28	88	1258
10	washingtondc	38.9	-77.0	0	1949	1949-03-29	88	1281

```
# i 247 more rows
```

```
# i 3 more variables: accumulative_growing_degree_days <dbl>,
```

```
# total_sunshine_duration <dbl>, total_precipitation <dbl>
```

Bayesian Regularized Neural Network

We did a 70/30 split of `complete_dataset` and stored the observations into a training set and `test_set` respectively. We trained the model using the training set and we will test how our model performs on unseen data with the test set. We will also use five fold cross-validation to tune any hyperparameters.

A bayesian regularized neural network was used to predict bloom_doy based on location, year, chill_hours, accumulative_growing_degree_days, total_sunshine_duration, and total_precipitation. The model with two neurons was chosen as the optimal model due to the low RMSE value achieved during cross-validation.

The test R-squared was approximately 0.884 which indicates that 88.4% of the variation in bloom dates can be explained by our model.

The test RMSE was about 2.7 which indicates that on average the bloom date predictions were off by 2.7 days.

```
# split data into a training set(70%) and a test set(30%)
set.seed(123)
index <- createDataPartition(y = complete_dataset$bloom_doy, p = .70, list = FALSE)
training_set <- complete_dataset[index, ]
test_set <- complete_dataset[-index, ]

# model will perform five fold cross-validation five times
five_fold_cross_validation <- trainControl(method = "repeatedcv", number = 5, repeats = 5)

brnn_model <- train(bloom_doy ~ location + year + chill_hours +
                    accumulative_growing_degree_days + total_sunshine_duration +
                    total_precipitation, data = training_set, method = "brnn",
                    trControl = five_fold_cross_validation)

print(brnn_model)
```

Bayesian Regularized Neural Networks

182 samples
6 predictor

No pre-processing

Resampling: Cross-Validated (5 fold, repeated 5 times)

Summary of sample sizes: 146, 146, 145, 145, 146, 146, ...

Resampling results across tuning parameters:

neurons	RMSE	Rsquared	MAE
---------	------	----------	-----

1	3.954411	0.8113908	3.201479
2	3.840946	0.8198470	3.075145
3	3.958073	0.8101347	3.148061

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was neurons = 2.

```
test_set_predictions <- predict(brnn_model, newdata = test_set)

test_set_performance <- postResample(pred = test_set_predictions, obs = test_set$bloom_doy

cat("\nTest Set Performance\n")
```

Test Set Performance

```
test_set_performance
```

	RMSE	Rsquared	MAE
	2.7024147	0.8844012	2.0381902

2024 Predictions

We used the same technique we used with `complete_dataset` to generate the daily and hourly weather aggregations for 2024 for each location. We then make predictions using our model trained on `complete_dataset` and output our predictions to a csv file.

```
brnn_model_all_data <- train(bloom_doy ~ location + year + chill_hours +
                             accumulative_growing_degree_days + total_sunshine_duration +
                             total_precipitation, data = complete_dataset, method = "brnn",
                             trControl = five_fold_cross_validation)
```

Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3728 alpha= 1.854 beta= 22.1614
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method

Scaling factor= 0.7023708
 gamma= 18.1144 alpha= 2.0888 beta= 26.2571
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037614
 gamma= 24.0938 alpha= 2.4621 beta= 27.4828
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3987 alpha= 1.9438 beta= 22.7664
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023479
 gamma= 17.9342 alpha= 2.3438 beta= 26.6186
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.703725
 gamma= 22.7067 alpha= 2.7762 beta= 26.841
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.429 alpha= 1.866 beta= 23.5163
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023708
 gamma= 18.474 alpha= 2.214 beta= 27.0853
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037614
 gamma= 24.2031 alpha= 2.1957 beta= 27.9005
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3166 alpha= 1.6751 beta= 19.9288
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023708
 gamma= 17.9334 alpha= 2.1029 beta= 23.4232
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037614
 gamma= 23.2342 alpha= 2.4919 beta= 23.638
 Number of parameters (weights and biases) to estimate: 11

Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3705 alpha= 1.6779 beta= 19.6764
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 16.6116 alpha= 2.3655 beta= 20.9063
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 20.9813 alpha= 2.4208 beta= 21.2811
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3849 alpha= 2.0358 beta= 23.0616
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023708
 gamma= 16.6242 alpha= 2.6332 beta= 25.3296
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037614
 gamma= 22.2348 alpha= 2.7472 beta= 26.0427
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.2644 alpha= 1.7561 beta= 21.6585
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.7155 alpha= 2.3504 beta= 24.1794
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 22.3694 alpha= 2.7348 beta= 24.382
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3969 alpha= 1.8382 beta= 20.088
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023708
 gamma= 18.2161 alpha= 1.5008 beta= 23.4191

Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037614
 gamma= 23.712 alpha= 2.2235 beta= 24.4075
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.4267 alpha= 1.7753 beta= 24.0614
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.7575 alpha= 2.364 beta= 26.9343
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 22.4748 alpha= 2.5832 beta= 27.2535
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3477 alpha= 1.7138 beta= 22.2494
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.9102 alpha= 2.1269 beta= 26.9819
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 22.8616 alpha= 2.4912 beta= 27.1781
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.4892 alpha= 1.7253 beta= 22.151
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.881 alpha= 2.0044 beta= 23.1886
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 24.2975 alpha= 1.8837 beta= 24.4613
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7

gamma= 10.3206 alpha= 1.8446 beta= 22.9951
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.2342 alpha= 2.3401 beta= 25.5086
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 22.6233 alpha= 2.4742 beta= 26.3865
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3115 alpha= 1.7504 beta= 22.4591
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023708
 gamma= 18.0597 alpha= 2.1556 beta= 26.4831
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037614
 gamma= 22.6076 alpha= 2.6723 beta= 27.0575
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3632 alpha= 2.0044 beta= 22.1233
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 16.9528 alpha= 2.376 beta= 25.6405
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 21.1642 alpha= 2.7326 beta= 25.6417
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.2922 alpha= 1.7832 beta= 20.8722
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023708
 gamma= 17.9548 alpha= 2.155 beta= 24.2104
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method

Scaling factor= 0.7037614
 gamma= 23.492 alpha= 2.3218 beta= 24.979
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3423 alpha= 1.7347 beta= 20.3603
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.5189 alpha= 2.2822 beta= 22.4593
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 22.8624 alpha= 2.38 beta= 23.0011
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3833 alpha= 1.8915 beta= 21.4346
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.2563 alpha= 2.1696 beta= 24.8405
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 23.9223 alpha= 2.2661 beta= 25.963
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3572 alpha= 1.8147 beta= 23.5775
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023708
 gamma= 17.821 alpha= 2.1901 beta= 28.094
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037614
 gamma= 23.0247 alpha= 2.6686 beta= 28.7677
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3775 alpha= 1.8395 beta= 23.6055
 Number of parameters (weights and biases) to estimate: 22

Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.7243 alpha= 2.2256 beta= 26.803
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 23.0467 alpha= 2.5506 beta= 27.241
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3534 alpha= 1.827 beta= 21.6844
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023708
 gamma= 17.4331 alpha= 2.3369 beta= 24.2312
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037614
 gamma= 22.3494 alpha= 2.6837 beta= 24.6244
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3589 alpha= 1.818 beta= 23.5589
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.7751 alpha= 2.2881 beta= 26.5833
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 22.4586 alpha= 2.6872 beta= 26.6256
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.4268 alpha= 1.8993 beta= 22.7685
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.9981 alpha= 2.2904 beta= 26.7874
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 23.6325 alpha= 2.5411 beta= 27.9419

Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3759 alpha= 1.7456 beta= 20.7865
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023708
 gamma= 17.3721 alpha= 2.2502 beta= 22.4702
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037614
 gamma= 21.939 alpha= 2.5026 beta= 22.7192
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3049 alpha= 1.8576 beta= 21.969
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023593
 gamma= 17.31 alpha= 2.3922 beta= 25.5431
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037431
 gamma= 21.5337 alpha= 2.7081 beta= 25.6131
 Number of parameters (weights and biases) to estimate: 11
 Nguyen-Widrow method
 Scaling factor= 0.7
 gamma= 10.3622 alpha= 1.7628 beta= 21.2871
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7023708
 gamma= 18.2868 alpha= 2.0831 beta= 24.8191
 Number of parameters (weights and biases) to estimate: 33
 Nguyen-Widrow method
 Scaling factor= 0.7037614
 gamma= 23.1982 alpha= 2.468 beta= 25.238
 Number of parameters (weights and biases) to estimate: 22
 Nguyen-Widrow method
 Scaling factor= 0.7018905
 gamma= 18.0881 alpha= 2.2963 beta= 25.7842


```

predictions_2024 <-
  tibble(
    location = c("washingtondc", "liestal", "kyoto", "vancouver", "newyorkcity"),
    year = 2024
  ) |>
  rowwise() |>
  mutate(
    bloom_date = Sys.Date() + 15,
    chill_hours = get_chill(location, year),
    accumulative_growing_degree_days = get_agdd(location, year, bloom_date),
    total_sunshine_duration = get_sunshine(location, year, bloom_date),
    total_precipitation = get_precip(location, year, bloom_date)
  )

predictions_2024 <- predictions_2024 |>
  dplyr::select(-bloom_date) |>
  bind_cols(predict(brnn_model_all_data, newdata = predictions_2024)) %>%
  rename(prediction = `...7`) |>
  mutate(prediction = round(prediction)) |>
  dplyr::select(location, prediction)

```

New names:

```
* `` -> `...7`
```

```
predictions_2024
```

```

# A tibble: 5 x 2
# Rowwise:
  location      prediction
  <chr>         <dbl>
1 washingtondc      77
2 liestal           76
3 kyoto             83
4 vancouver         84
5 newyorkcity       76

```

```
write.csv(predictions_2024, file = "cherry-predictions.csv", row.names = FALSE)
```